

## Online Resource 2

# Datasets, Performance Comparisons, Visualizations, and Evaluation Metrics for the Nova Optimizer

Ali Zeydi Abdian<sup>1</sup>, Mohammad Masoud Javidi<sup>2</sup> and Najme Mansouri<sup>3</sup>

<sup>1</sup>Department of Computer Science, Shahid Bahonar University of Kerman, Kerman, Iran, E-mail: ali.zeydi.abdian@ut.ac.ir, alizeydiabdian@math.uk.ac.ir  
(Corresponding Author)

<sup>2</sup> Professor, Faculty of Shahid Bahonar University of Kerman, Kerman, Iran, Email: javidi@uk.ac.ir

<sup>3</sup> Associate Professor, Faculty of Shahid Bahonar University of Kerman, Kerman, Iran, Email: n.mansouri@uk.ac.ir, najme.mansouri@gmail.com, Box No. 76135-133, Fax: 03412111865

Corresponding author: alizeydiabdian@math.uk.ac.ir

**Journal Name: Signal, Image and Vide Processing**

## Contents

<b>1</b>	<b>Datasets</b>	<b>2</b>
<b>2</b>	<b>Time and Space Complexity of Common First-Order Optimizers</b>	<b>2</b>
<b>3</b>	<b>Evaluating the Performance of Various Optimizers on Some Well-Known Networks: A Comparison with Nova</b>	<b>3</b>
<b>4</b>	<b>Confusion Matrix</b>	<b>6</b>
<b>5</b>	<b>Results</b>	<b>7</b>
5.1	Comparison of Performance Metrics of Nova with Other Optimizers . . . . .	7
5.2	Learning Curves and Bar Plots . . . . .	9
5.3	Confusion Matrices of datasets of CIFAR-10, MNIST, Noisy MNIST and SST-2 . . . . .	11
5.4	ROC and AUC . . . . .	15

# 1 Datasets



Fig 1. Visualization of CIFAR-10 and MNIST datasets.




Fig 2. MNIST and noisy MNIST

## 2 Time and Space Complexity of Common First-Order Optimizers

Let  $P$  denote the total number of trainable parameters in a model. Each optimizer updates all parameters in one pass, performing a constant number of element-wise operations per parameter, so the *time complexity* per iteration is  $\mathcal{O}(P)$ .

Many optimizers maintain auxiliary state vectors of length  $P$ . If an optimizer keeps  $k$  such vectors, its

additional memory complexity is:

$$\mathcal{O}(kP) \quad (\text{beyond the } P \text{ parameters and their gradients}).$$

Table 1 summarizes the relative computational efficiency and memory usage of various optimizers

Table 1. Comparison of time and space complexity for popular gradient-based optimizers. Here  $P$  is the number of model parameters; “Memory overhead” counts only the auxiliary state beyond parameters and gradients.

Optimizer	Time per update	Memory overhead	State vectors per parameter
SGD	$\mathcal{O}(P)$	$\mathcal{O}(0 \cdot P)$	None
SGD + Momentum	$\mathcal{O}(P)$	$\mathcal{O}(1 \cdot P)$	1 (velocity)
RMSprop	$\mathcal{O}(P)$	$\mathcal{O}(1 \cdot P)$	1 (running $\mathbb{E}[g^2]$ )
Adam / AdamW	$\mathcal{O}(P)$	$\mathcal{O}(2 \cdot P)$	2 (first and second moments: $m, v$ )
NAdam	$\mathcal{O}(P)$	$\mathcal{O}(2 \cdot P)$	2 (as in Adam)
Nova (AMSGrad-style)	$\mathcal{O}(P)$	$\mathcal{O}(3 \cdot P)$	3 ( $m, v$ , and $\max(v)$ )

### 3 Evaluating the Performance of Various Optimizers on Some Well-Known Networks: A Comparison with Nova

Table 2. ResNet-18 architecture details

Layer / Block	Operation	Output / Notes
Input	RGB Image	$224 \times 224 \times 3$
Conv1	$7 \times 7$ Convolution, 64 filters, stride 2, padding 3	$112 \times 112 \times 64$
BN + ReLU	—	—
MaxPool	$3 \times 3$ Max Pooling, stride 2	$56 \times 56 \times 64$
<b>Residual Blocks (each block: 2 conv layers, BN, ReLU)</b>		
Layer 1 (conv2_x)	2 blocks with 64 filters (stride 1)	$56 \times 56 \times 64$
Layer 2 (conv3_x)	2 blocks with 128 filters; first block uses stride 2	$28 \times 28 \times 128$
Layer 3 (conv4_x)	2 blocks with 256 filters; first block uses stride 2	$14 \times 14 \times 256$
Layer 4 (conv5_x)	2 blocks with 512 filters; first block uses stride 2	$7 \times 7 \times 512$
Global Average Pooling	Pooling over $7 \times 7$ spatial dimensions	$1 \times 1 \times 512$
FC Layer	Fully Connected, e.g. 1000 classes for ImageNet	1000 (or as specified)

Table 3. LeNet-5 architecture details

Layer	Operation	Output / Notes
Input	Grayscale image	$32 \times 32 \times 1$
C1	Convolution ( $5 \times 5$ kernel), 6 filters, activation tanh	$28 \times 28 \times 6$
S2	Average pooling (subsampling)	$14 \times 14 \times 6$
C3	Convolution ( $5 \times 5$ kernel), 16 filters, tanh	$10 \times 10 \times 16$
S4	Average pooling (subsampling)	$5 \times 5 \times 16$
C5	Convolution (fully-connected), kernel size matching input, tanh	$1 \times 1 \times 120$
F6	Fully connected	84 units
Output	Fully connected	10 units (for digit classification)

Table 4. Comparison of ResNet-18 and LeNet-5

Attribute	ResNet-18	LeNet-5
Year developed	2015	1998
Depth (Layers)	18	7
Key Innovation	Residual (skip) connections	Convolution + pooling
Original application	Large-scale image recognition	Handwritten digit recognition
Training techniques	Batch normalization, deeper architecture	Simpler activations (e.g., tanh)
Network complexity	High (designed for complex datasets)	Low (optimized for limited variability)

Table 5. Used architecture CIFAR-10 ResNet-18 in the paper

Layer/Block	Modification	Parameters/Notes
Input	CIFAR-10 Image	32×32×3
Conv1	3×3 Conv, stride 1	Original: 7×7 stride 2
MaxPool	Removed	Identity operation
Residual Blocks	Added Spatial Dropout	Layer1: 0.05, Layer2: 0.1 Layer3: 0.15, Layer4: 0.2
Classifier	Modified FC Layers Added Dropout	512 → 256 → 10 FC dropout: 0.3

Table 6. Used architecture MNIST LeNet-5 in the paper

Layer	Operation	Parameters
Input	MNIST Image	28×28×1
Conv1	5×5 Conv, 6 filters	Output: 24×24×6
MaxPool	2×2	Output: 12×12×6
Conv2	5×5 Conv, 16 filters	Output: 8×8×16
MaxPool	2×2	Output: 4×4×16
FC1	Linear Layer	256 → 120
FC2	Linear Layer	120 → 84
FC3	Linear Layer	84 → 10
Regularization	Dropout	p=0.25 after FC1

Table 7. Training configuration

Category	Parameter	Value
Training	Epochs	30
	Batch Size	128
	LR Scheduler	StepLR (gamma=0.9)
Optimizers	Base LR	0.001
	Weight Decay	0.01
	Nova	betas=(0.9, 0.999), eps=1e-8
	SGD	momentum=0.9
Data Augmentation	CIFAR-10	RandomCrop(32,4)+Flip
	MNIST	None
Normalization	CIFAR-10	Mean=[0.4914, 0.4822, 0.4465]
		Std=[0.2023, 0.1994, 0.2010]
	MNIST	Mean=0.1307, Std=0.3081

Table 8. Training configuration for SST-2

Category	Parameter	Value
Training	Epochs	10
	Batch Size	64
	LR Scheduler	StepLR (gamma=0.9)
Optimizers	Base LR	0.001
	Weight Decay	0.0
	Nova	betas=(0.9, 0.999), eps=1e-8
	Adam	Default parameters
Data Augmentation	SST-2	None
Normalization	Preprocessing	Lowercase conversion, tokenization, and padding
	Vocabulary Construction	Count-based, with '<unk>' and '<pad>' tokens

Table 9. Detailed architecture of SST-2 text classification model

Component	Description
Model Type	Feedforward Neural Network (Embedding + Linear Classifier)
Embedding Layer	<ul style="list-style-type: none"> <li>Input dimension: vocab_size (determined from training data)</li> <li>Output dimension (embed_dim): 100 (configurable)</li> <li>Special tokens: &lt;unk&gt;=0, &lt;pad&gt;=1 (for OOV words and padding)</li> <li>Trainable parameters: vocab_size × embed_dim</li> </ul>
Pooling	<ul style="list-style-type: none"> <li>Operation: Mean pooling across sequence length (dim=1)</li> <li>Reduces variable-length sequences to fixed-size embedding</li> </ul>
Classifier Layer	<ul style="list-style-type: none"> <li>Single Linear layer (embed_dim → num_class)</li> <li>Output size: 2 (binary classification: Negative/Positive)</li> <li>No explicit activation (CrossEntropyLoss includes LogSoftmax)</li> <li>Trainable parameters: embed_dim × num_class + bias</li> </ul>
Forward Pass	<ol style="list-style-type: none"> <li>1. Tokenize text → convert to vocab indices</li> <li>2. Embedding layer: maps indices to dense vectors</li> <li>3. Mean pooling: aggregates sequence embeddings</li> <li>4. Linear layer: produces logits for classification</li> </ol>
Training Hyperparameters	<ul style="list-style-type: none"> <li>Batch size: 64</li> <li>Loss: CrossEntropyLoss (combines LogSoftmax + NLLLoss)</li> <li>Optimizer: Configurable (e.g., Adam, SGD)</li> <li>Learning rate scheduler: StepLR (step_size=1, gamma=0.9)</li> </ul>
Data Pipeline	<ul style="list-style-type: none"> <li>Dataset: SST-2 (Stanford Sentiment Treebank v2)</li> <li>Splits: 90% train / 10% val (from original train), test=validation set</li> <li>Preprocessing: Lowercase + NLTK word_tokenize</li> <li>Vocabulary: Built from training set, with &lt;unk&gt; and &lt;pad&gt;</li> <li>Padding: Dynamic per batch (to max length in batch)</li> </ul>

## 4 Confusion Matrix

Table 10. Confusion matrix and evaluation metrics

		Predicted class	
		Positive	Negative
Actual class	Positive	True Positive (TP)	False Negative (FN)
	Negative	False Positive (FP)	True Negative (TN)
		Precision = $\frac{TP}{TP+FP}$	Recall = $\frac{TP}{TP+FN}$
		Accuracy = $\frac{TN+TP}{TP+TN+FP+FN}$	F1 score = $\frac{2TP}{2TP+FP+FN}$

## 5 Results

### 5.1 Comparison of Performance Metrics of Nova with Other Optimizers

Table 11. Performance metrics on CIFAR-10

Optimizer	Train Loss	Train Acc (%)	Val Loss	Val Acc (%)	Test Loss	Test Acc (%)	Epoch Time (s)
Nova	0.1550	94.70	0.3180	89.42	0.3345	90.01	41.95
Adam	0.4788	84.94	0.5094	83.08	0.5069	83.14	39.99
RMSprop	1.1128	59.39	1.0714	60.36	1.0597	61.01	38.50
SGD	1.2252	55.11	1.1721	57.88	1.1325	58.57	39.87
Nadam	0.5583	82.09	0.5642	81.20	0.5718	80.85	40.21

Table 12. Classification metrics on CIFAR-10

Optimizer	Precision (%)	Recall (%)	F1-score (%)
Nova	89.98	90.01	89.99
Adam	83.21	83.14	83.00
RMSprop	61.01	61.01	59.77
SGD	58.89	58.57	58.18
Nadam	80.93	80.85	80.82

Table 13. Performance metrics on MNIST

Optimizer	Train Loss	Train Acc (%)	Val Loss	Val Acc (%)	Test Loss	Test Acc (%)	Epoch Time (s)
Nova	0.0204	99.33	0.0370	98.86	0.0295	98.95	12.01
Adam	0.0673	98.20	0.0578	98.10	0.0504	98.58	11.22
RMSprop	0.0699	98.08	0.0631	97.94	0.0514	98.50	11.21
SGD	0.1240	96.45	0.0960	97.12	0.0855	97.51	11.16
Nadam	0.0661	98.21	0.0582	98.16	0.0484	98.66	11.34

Table 14. Classification metrics on MNIST

Optimizer	Precision (%)	Recall (%)	F1-score (%)
Nova	98.95	98.94	98.94
Adam	98.58	98.57	98.57
RMSprop	98.50	98.48	98.49
SGD	97.52	97.49	97.50
Nadam	98.67	98.65	98.66

Table 15. Performance metrics for SST-2

Optimizer	Test Loss	Test Acc (%)	Precision (%)	Recall (%)	F1-score (%)	Epoch Time (s)
Nova	0.5125	82.00	82.00	82.01	81.99	15.48
Adam	0.6397	66.28	68.71	65.94	64.85	15.31
AdamW	0.5804	81.65	81.69	81.62	81.63	15.28
RMS	0.6422	65.94	69.05	65.55	64.15	15.23
SGD	0.6810	54.47	56.64	53.88	48.85	15.08
Nadam	0.6399	66.17	68.84	65.81	64.61	15.36

Table 16. Classification metrics on SST-2

Optimizer	Precision (%)	Recall (%)	F1-score (%)
Nova	82.11	82.12	82.11
Adam	69.56	66.38	65.19
AdamW	80.93	80.39	80.39
RMS	68.10	65.48	64.41
SGD	59.31	53.28	44.63
Nadam	69.62	66.13	64.81

Table 17. Performance metrics for noisy MNIST

Optimizer	Test Loss	Test Acc (%)	Precision (%)	Recall (%)	F1-score (%)	Epoch Time (s)
Nova	0.1065	96.58	96.56	96.56	96.56	8.47
AdamW	0.1112	96.14	96.14	96.11	96.12	7.85
Adam	0.1156	96.17	96.16	96.15	96.14	7.77
RMS	0.1072	96.53	96.52	96.50	96.50	7.75
SGD	0.2727	91.32	91.24	91.21	91.20	7.70
Nadam	0.1052	96.43	96.42	96.39	96.40	7.63

Table 18. Performance of Nova optimizer hyperparameter combinations on MNIST

Combination	LR	$\beta_1$	$\beta_2$	WD	Test Acc (%)	Precision (%)	Recall (%)	F1-score (%)
Comb1	0.001	0.9	0.999	0.01	99.21	99.21	99.20	99.20
Comb2	0.001	0.9	0.99	0.01	99.17	99.16	99.16	99.16
Comb3	0.001	0.95	0.999	0.001	99.11	99.11	99.10	99.11
Comb4	0.001	0.95	0.99	0.001	98.94	98.94	98.93	98.93
Comb5	0.01	0.9	0.999	0.01	99.14	99.14	99.13	99.13
Comb6	0.01	0.9	0.99	0.01	99.01	99.01	98.99	99.00
Comb7	0.01	0.95	0.999	0.001	98.44	98.44	98.42	98.43
Comb8	0.01	0.95	0.99	0.001	99.00	99.00	98.99	98.99



## 5.2 Learning Curves and Bar Plots

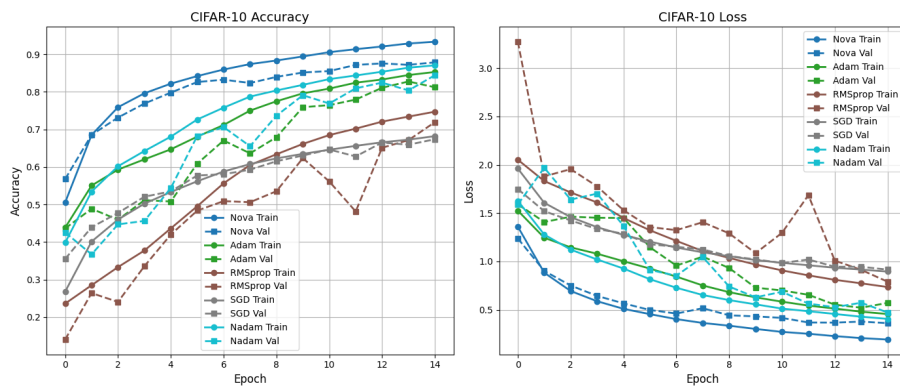


Fig 3. Learning curve CIFAR-10 with different optimizers

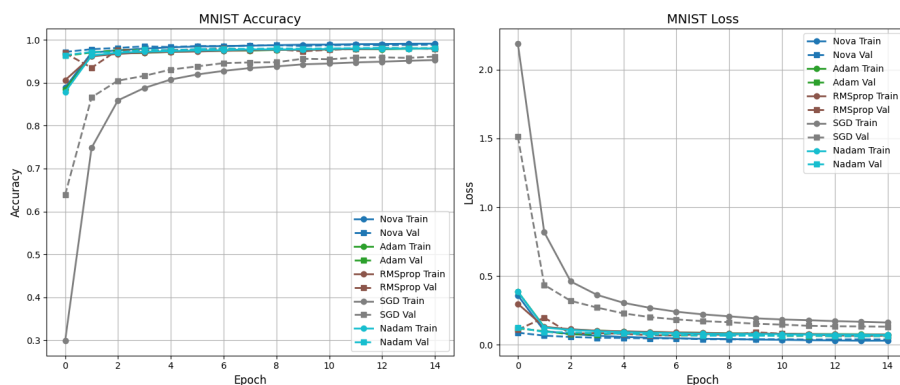


Fig 4. Learning curve MNIST with different optimizers

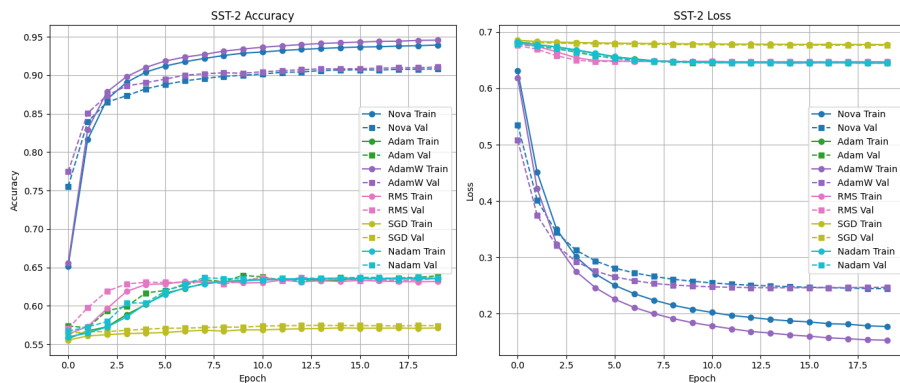


Fig 5. Learning curve SST-2 with different optimizers

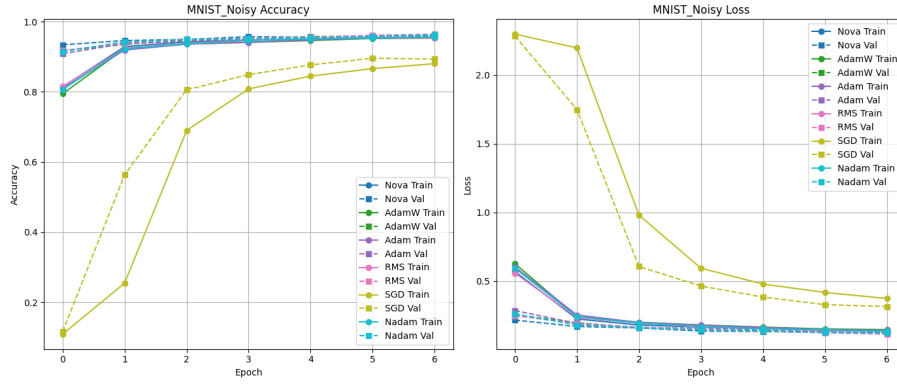


Fig 6. Learning curve Noisy MNIST with different optimizers





Fig 7. Bar plot for comparing accuracy of Nova with other optimizers (CIFAR-10 and MNIST)

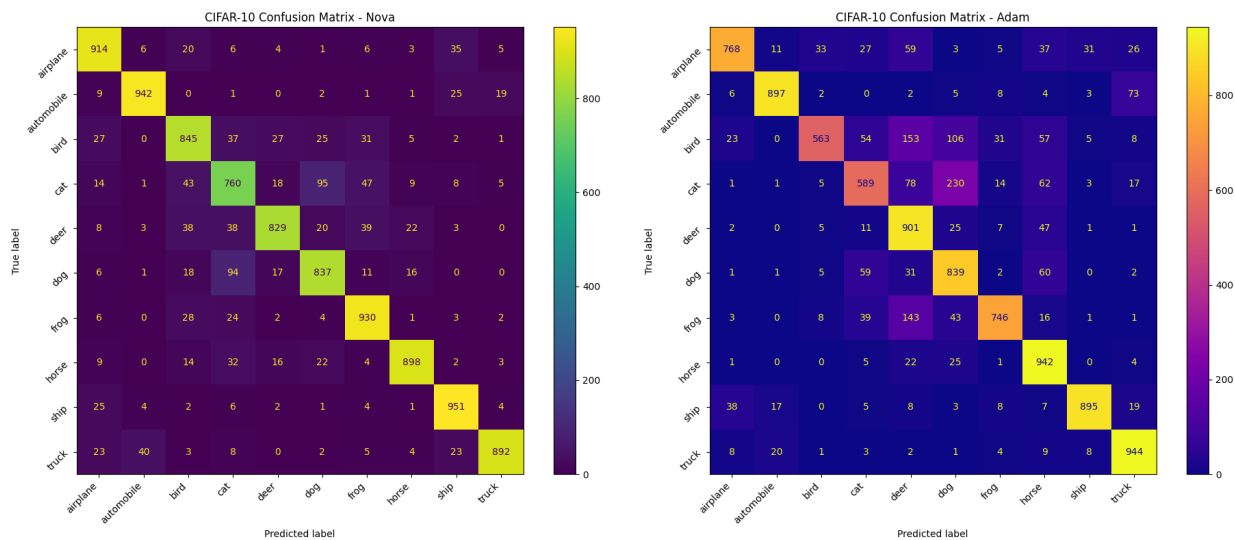


(a) Bar plot for comparing accuracy of Nova with other optimizers (SST-2)

(b) Bar plot for comparing accuracy of Nova with other optimizers (noisy MNIST)


Fig 8. Bar plot for comparing accuracy of Nova with other optimizers

5.3 Confusion Matrices of datasets of CIFAR-10, MNIST, Noisy MNIST and SST-2




(a) Confusion matrix of CIFAR-10 (Nova)

(b) Confusion matrix of CIFAR-10 (Adam)

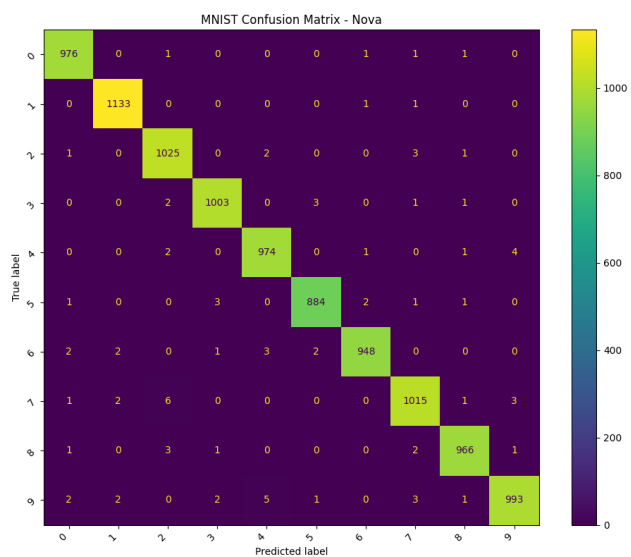


(c) Confusion matrix of CIFAR-10 (RMSprop)

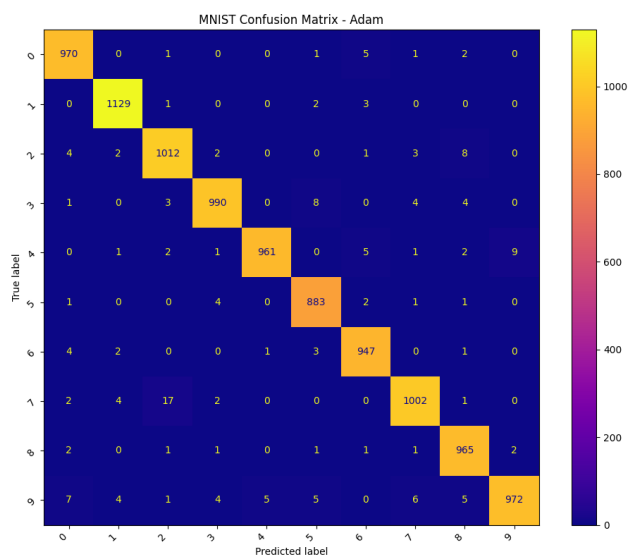
(d) Confusion matrix of CIFAR-10 (SGD)



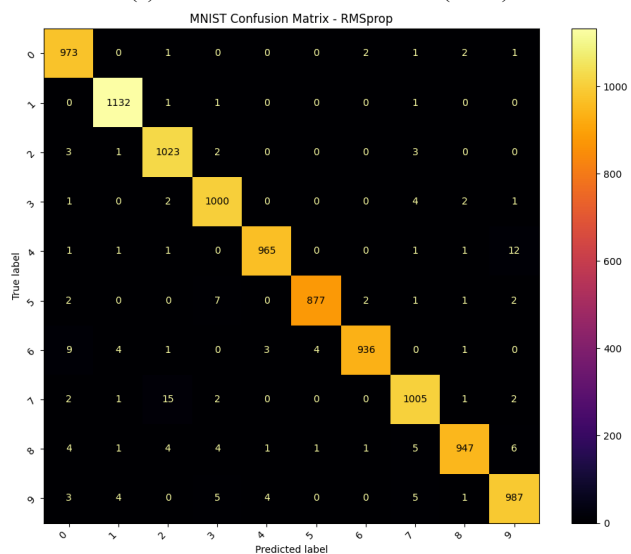
(e) Confusion matrix of CIFAR-10 (Nadam)



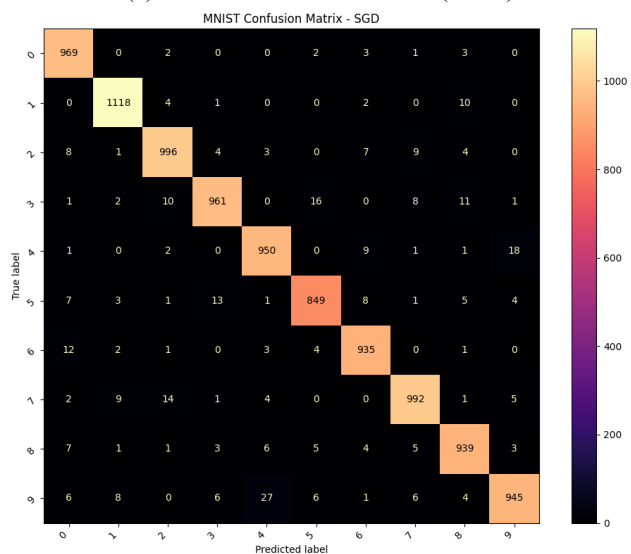
(a) Confusion matrix of MNIST (Nova)




(b) Confusion matrix of MNIST (Adam)



(c) Confusion matrix of MNIST (RMSprop)

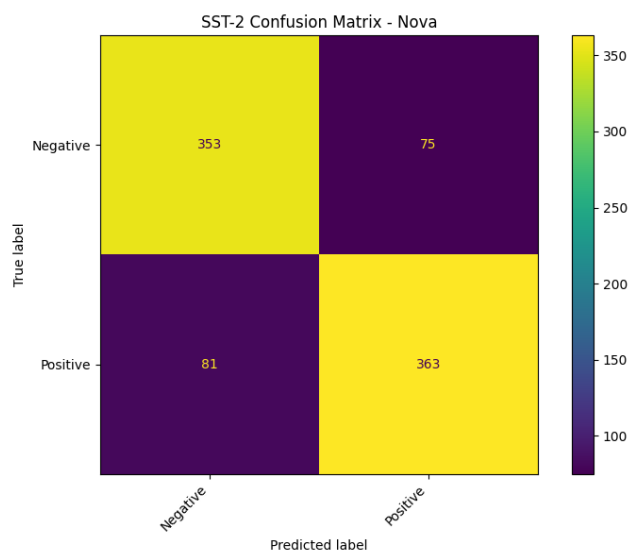


(d) Confusion matrix of MNIST (SGD)




(e) Confusion matrix of MNIST (Nadam)


Fig 10. Confusion matrix for dataset with respect to different optimizers



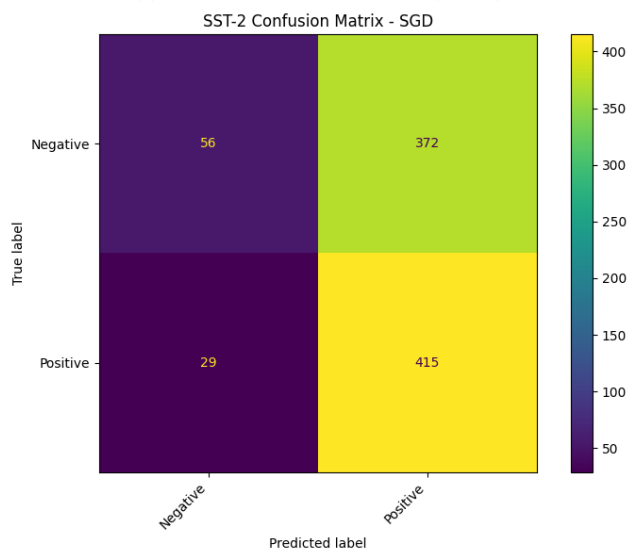
(a) Confusion matrix of SST-2 (Nova)




(b) Confusion matrix of SST-2 (Adam)




(c) Confusion matrix of SST-2 (RMSprop)



(d) Confusion matrix of SST-2 (SGD)

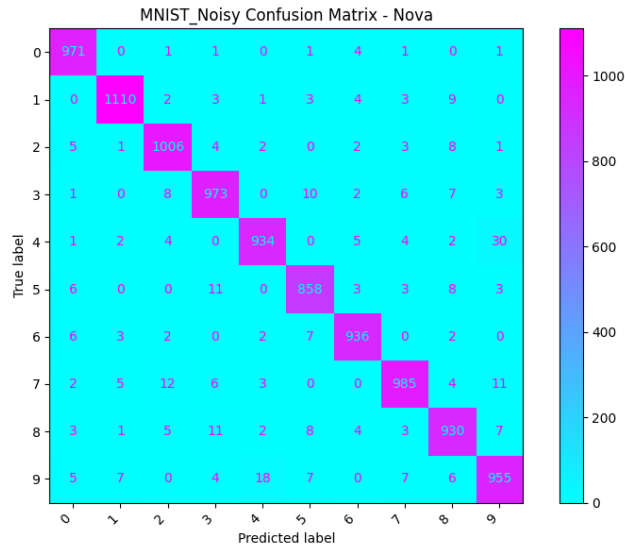


(e) Confusion matrix of SST-2 (Nadam)

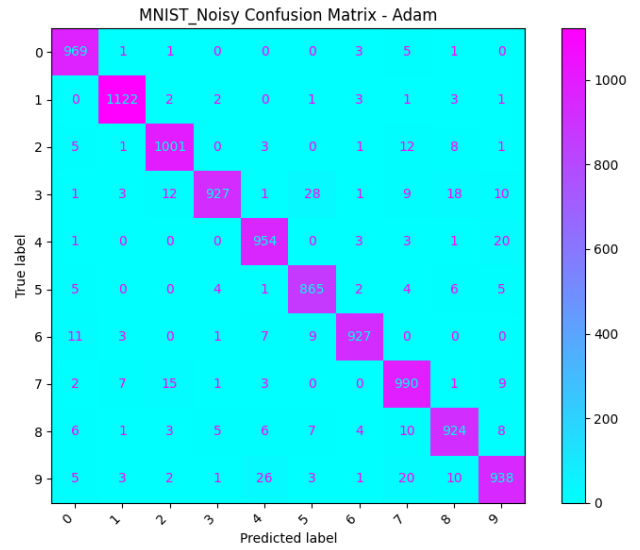


(f) Confusion matrix of SST-2 (AdamW)

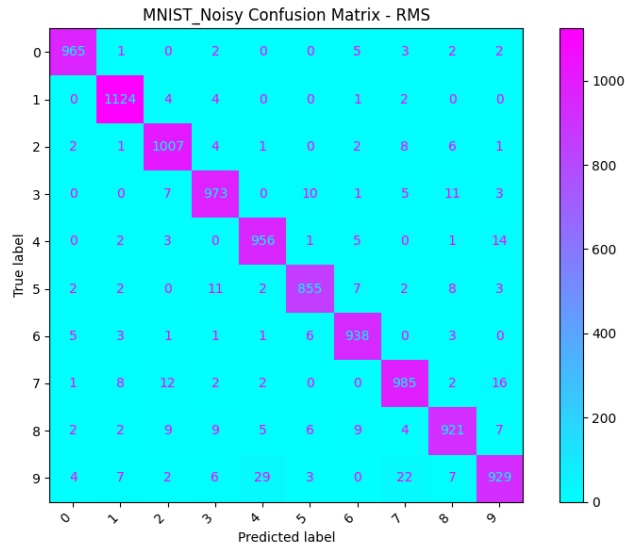
Fig 11. Confusion matrix for dataset with respect to different optimizers




(a) Confusion matrix of noisy MNIST (Nova)




(b) Confusion matrix of noisy MNIST (Adam)




(c) Confusion matrix of noisy MNIST (RMSprop)



(d) Confusion matrix of noisy MNIST (SGD)



(e) Confusion matrix of noisy MNIST (Nadam)



(f) Confusion matrix of noisy MNIST (AdamW)

Fig 12. Confusion matrix for dataset with respect to different optimizers

## 5.4 ROC and AUC

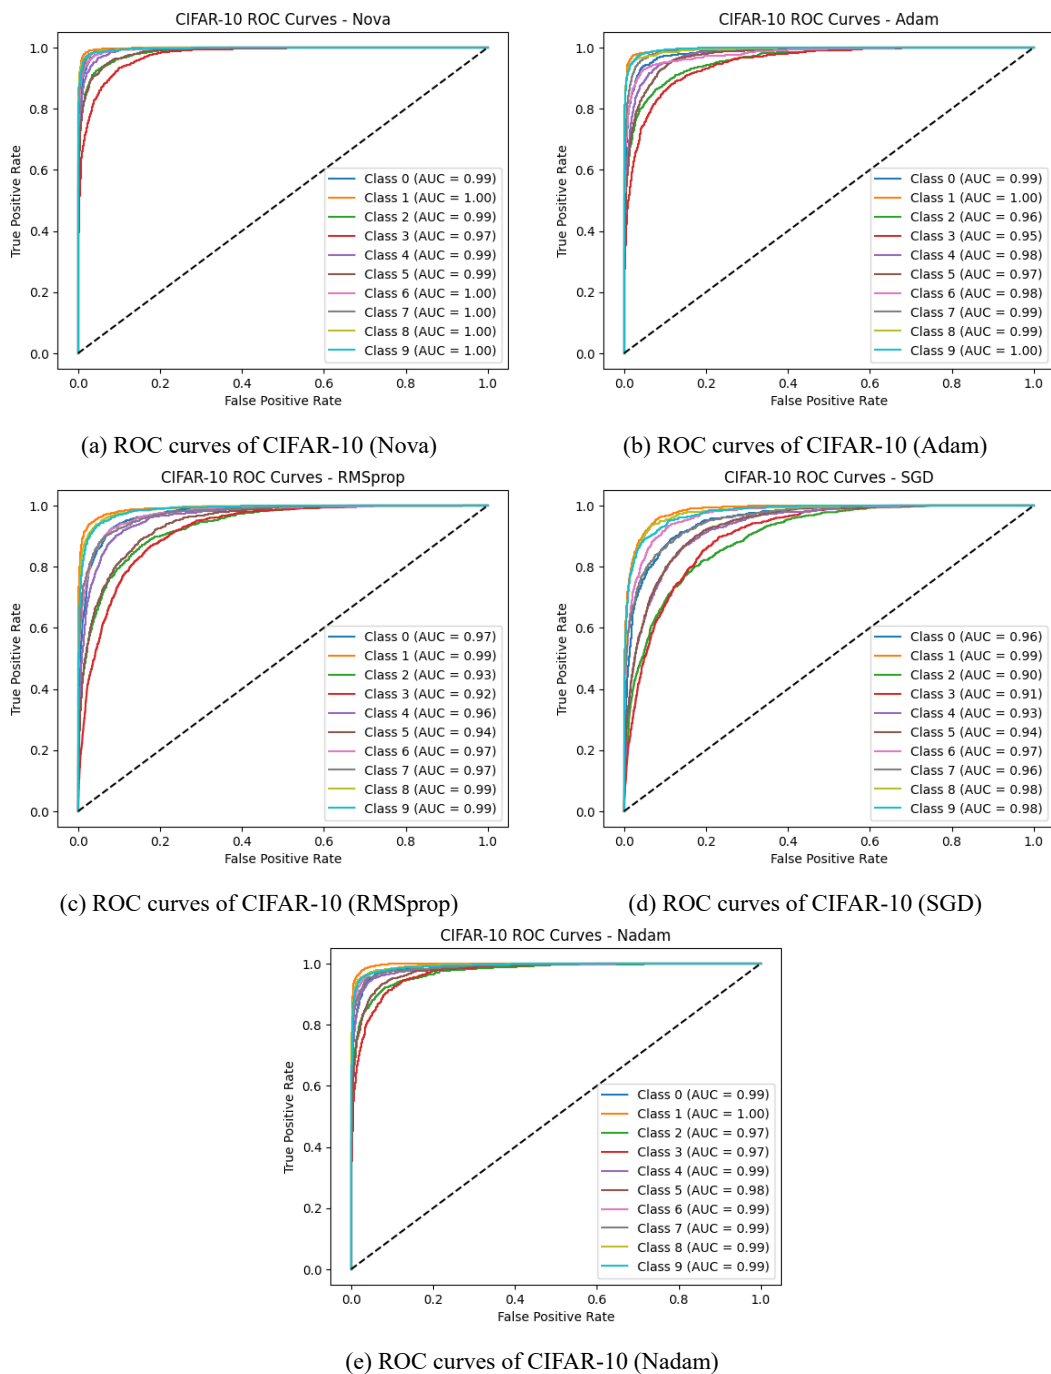
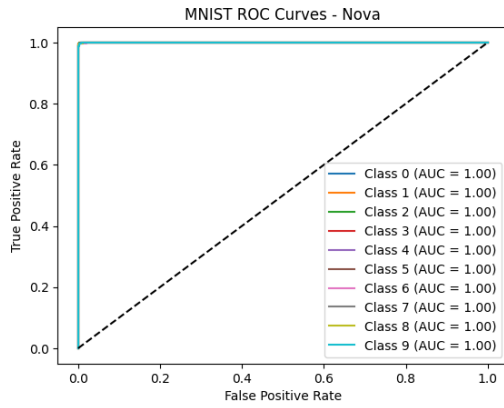
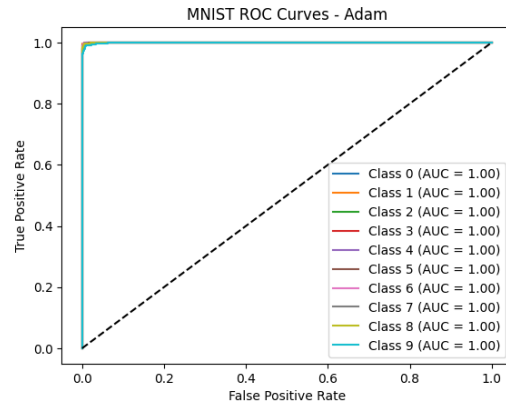


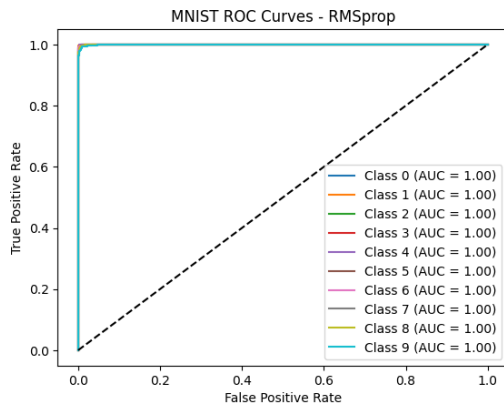
Fig 13. ROC curves for CIFAR-10 dataset with respect to different optimizers



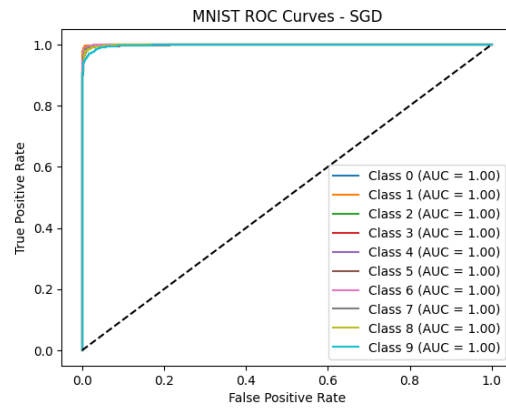
(a) ROC curves of MNIST (Nova)



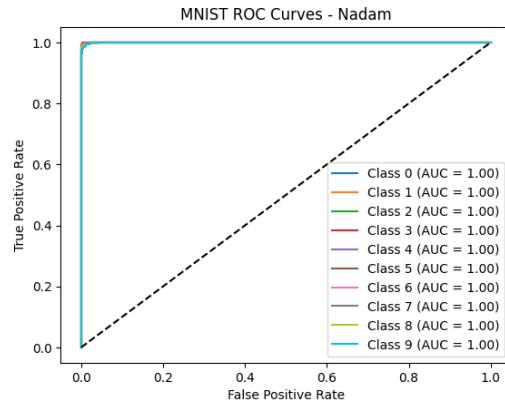
(b) ROC curves of MNIST (Adam)



(c) ROC curves of MNIST (RMSprop)



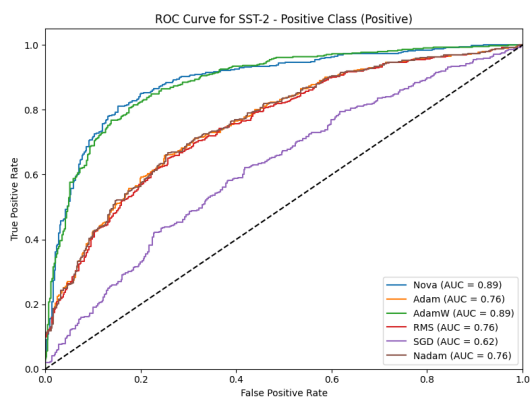
(d) ROC curves of MNIST (SGD)



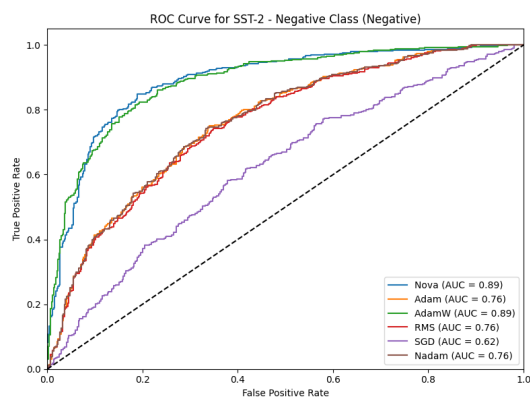
(e) ROC curves of MNIST (Nadam)

Fig 14. ROC curves for MNIST dataset with respect to different optimizers





(a) ROC of SST-2 (Positive class)



(b) ROC of SST-2 (Negative class)

Fig 15. ROC curves for SST-2 dataset with respect to different optimizers